



Differentiable Inductive Logic Programming (∂ ILP) for Fraud Detection

Data Science Master's Thesis - 2024

Boris Wolfson, Dr. Erman Acar [†]

September 5, 2024

University of Amsterdam

[†]Institute for Logic, Language and Computation

CODE

Github repo: <https://github.com/wolfbrr/ThesisDS>



OUTLINE

1. Introduction

2. ∂ ILP

3. Framework

Introduction

INTRODUCTION: LOGIC PROGRAMMING - DEFINITIONS

Logic programming

- Family of programming languages in which the central component is an if-then rule/clause
- The rules of an ILP framework are written as a set of definite clauses of the following form:

$$\alpha \leftarrow \alpha_1, \alpha_2, \dots, \alpha_n, \tag{1}$$

Here:

- A set of atoms $\{\alpha_i\}$ defines a body of the rule.
- α is the head atom of the rule, and equals to True, if all the atoms in the body are True
- Each atom is defined by a predicate

INTRODUCTION: LOGIC PROGRAMMING-EXAMPLES

Example - predicates:

- $Q(X,Y)$: Q is a predicate of arity 2, can be "x is greater than y" and "x is the father of y", or x transferred money to y

Example - Logic Program

- The following program defines the set of rules \mathcal{R} for connected relations as the transitive closure of the edge relation:

$$\begin{aligned} \text{connected}(X, Y) &\leftarrow \text{edge}(X, Y) \\ \text{connected}(X, Y) &\leftarrow \text{edge}(X, Z), \text{connected}(Z, Y). \end{aligned} \tag{2}$$

INDUCTIVE LOGIC PROGRAMMING

Idea: Based on the background knowledge \mathcal{B} , to derive a set of rules \mathcal{R} for a target predicate such that it [Muggleton, 1991]:

- entails all the positive examples \mathcal{P}
- does not entail the negative examples \mathcal{N}

INDUCTIVE LOGIC PROGRAMMING-EXAMPLE

Example - Even numbers [Evans and Grefenstette, 2018], learn target predicate $even(X)$:

- $\mathcal{B} = \{zero(0), succ(0, 1), succ(1, 2), succ(2, 3), succ(3, 4), succ(4, 5)\}$
- $\mathcal{P} = \{even(0), even(2), even(4)\}$
- $\mathcal{P} = \{even(1), even(3), even(5)\}$

Possible Solution

$$\begin{aligned} even(X) &\leftarrow zero(X) \\ even(X) &\leftarrow even(Y), succ2(Y, X) \\ succ2(X, Y) &\leftarrow succ(X, Z), succ(Z, Y) \end{aligned} \tag{3}$$

∂ ILP

∂ ILP

∂ ILP:

- Combines advantages of ILP and the neural-network-based systems - Neurosymbolic (NeSy) AI.
- Has two types of predicates:
 - The intensional P_i , target and auxiliary predicates $\{p_a\}$,
 - The extensional P_e , from the background knowledge ($zero(X)$, $edge(X, Y)$, ...)
- Generates a list of possible definite clauses for intentional predicates based on the program template [Tausend, 1994]

∂ILP - PROGRAM TEMPLATE

Each intensional predicate p_a is defined by its arity, and is described by two clauses according to rule templates $(\tau_{p_a}^1, \tau_{p_a}^2)$.

$$\tau = (n_{\exists}, int) \tag{4}$$

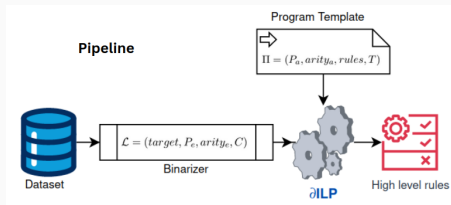
- n_{\exists} - the number of existentially quantified variables allowed in the clause
- int is a flag that determines whether the atoms in the clause can use intensional predicates P_i

∂ILP learns a set of clauses to minimize the loss function.

Framework

PIPELINE

- *Dataset* - Paysim [Lopez-Rojas et al., 2016], and synthetically generated scenarios
- *Binarizer* converts numerical columns to binary and creates sets of facts, positive and negative examples.
- *High level rules* generated SQL query from the derived rules



RESULTS

Chain of fraud


- Transaction(X,Y) - X sends transaction to Y
- Fraud(Z,X) - Transaction from Z to X is Fraud
- Fraud_Chain(X,Y) - X and Y are in a chain of a fraud event
- Template $\tau_{target}^1 = (n_3 = 1, int = 0)$
- Rule $Fraud_Chain(X,Y) \leftarrow Fraud(Z,X), Transaction(X,Y)$



	orig	destination	Fraud_chain-orig-destination
0	16058	16005	False
1	16005	16052	False
2	16039	16007	False
3	16086	16014	True
4	16043	16004	False
5	16011	16007	True
6	16002	16011	False
7	16051	16086	False
8	16080	16077	False

Fraud Relationship

- Fraudsters(X,Y) - X and Y are Fraudsters
- Fraud(X,Y) - Transaction from X to Y is Fraud
- Template $\tau_{target}^1 = (n_3 = 0, int = 1)$
 $\tau_{target}^2 = (n_3 = 1, int = 1)$
- Rule



	X	Y	Fraudsters-X-Y	Fraud-X-Y
0	1	2	True	True
1	2	3	True	True
2	3	4	True	True
3	2	1	True	True
4	1	3	True	False
5	1	4	True	False
6	1	1	True	False
7	2	4	True	False
8	2	2	True	False

PaySim Dataset

- Binarised Columns from Decision Trees Thresholds
- X - transaction id
- Template

$$\begin{aligned}\tau_{target}^1 &= (n_3 = 0, int = 1) \\ \tau_{target}^2 &= (n_3 = 0, int = 1) \\ \tau_{p1}^1 &= (n_3 = 0, int = 1) \\ \tau_{p1}^2 &= None \\ \tau_{p2}^1 &= (n_3 = 0, int = 0) \\ \tau_{p2}^2 &= None\end{aligned}$$

- Rule

isFraud(X_0) \leftarrow NOT (oldbalanceDest > -0.007)(X_0), pred2(X_0)
isFraud(X_0) \leftarrow pred2(X_0), amount > 1.297(X_0)
pred1(X_0) \leftarrow NOT (oldbalanceDest > -0.007)(X_0), pred2(X_0)
pred2(X_0) \leftarrow external_dest(X_0), type_TRANSFER(X_0)

Performance	oILP	DT	DSC
Accuracy	0.999	0.999	0.999
Precision	0.973	0.971	0.984
Recall	0.501	0.665	0.501
F1	0.662	0.789	0.664
MCC	0.698	0.803	0.702



PROS/CONS

- Pros:
 - ∂ ILP generalizes from a small amount of data - Paysim results based on 1000 from the database out of 6 million transactions
 - Can create recursion predicates
 - Provided shorter explainable rules than the Decision Tree approach
- Cons:
 - Scalability
 - Did not outperform other techniques - showed low recall
 - Expert knowledge - Program Template definition

REFERENCES I



Evans, R. and Grefenstette, E. (2018).

Learning explanatory rules from noisy data.

Journal of Artificial Intelligence Research, 61:1–64.



Lopez-Rojas, E., Elmir, A., and Axelsson, S. (2016).

Paysim: A financial mobile money simulator for fraud detection.

In *28th European Modeling and Simulation Symposium, EMSS, Larnaca*, pages 249–255. Dime University of Genoa.



Muggleton, S. (1991).

Inductive logic programming.

New generation computing, 8:295–318.

REFERENCES II



Tausend, B. (1994).

Representing biases for inductive logic programming.

In *European Conference on Machine Learning*, pages 427–430.

Springer.

POSTER

Differentiable Inductive Logic Programming (∂ILP) for Fraud Detection

Boris Wolfson, Dr. Erman Acar (Data Science Master's Thesis)



UNIVERSITEIT VAN AMSTERDAM

∂ILP

- "Learning Explanatory Rules from Noisy Data", Richard Evans, Edward Grefenstette
- Learns set of rules $\alpha \leftarrow \alpha_1, \dots, \alpha_m$
- Required input: Positive, Negative examples and a set of facts
- Neurosymbolic AI - Deep NN + Symbolic Reasoning

Objectives:

- Compare performance to classical rule generation methods: Deep Symbolic Regression, Decision Tree
- Test Recursive Structures for the fraudulent relationships detection

Contributions

- Pipeline for converting tabular dataset to required input format
- Set of templates for different programs

Examples

Chain of fraud

- Transaction(X,Y) - X sends transaction to Y
- Fraud(Z,X) - Transaction from Z to X is Fraud
- Fraud_Chain(X,Y) - X and Y are in a chain of a fraud event
- Template $\tau_{\text{target}}^{\text{Chain}} = \{R_3 = 1, \text{int} = 0\}$
- Rule $\text{Fraud_Chain}(X,Y) \leftarrow \text{Fraud}(Z,X), \text{Transaction}(X,Y)$



city	destination	Fraud_Chain-city-destination
0	10000	10000
1	10000	10000
2	10000	10000
3	10000	10000
4	10000	10000
5	10000	10000
6	10000	10000
7	10000	10000
8	10000	10000

Fraud Relationship

- Fraudsters(X,Y) - X and Y are Fraudsters
- Fraud(X,Y) - Transaction from X to Y is Fraud
- Template $\tau_{\text{target}}^{\text{Fraud}} = \{R_3 = 0, \text{int} = 1\}$
- Rule $\text{Fraudsters}(X,Y) \leftarrow \text{Fraud}(X,Y)$




X	Y	Fraudsters-X-Y
0	1	True
1	2	True
2	3	True
3	4	True
4	5	True
5	6	True
6	7	True
7	8	True
8	9	True

PaySim Dataset

- Binarised Columns from Decision Trees Thresholds
- X - transaction id
- Template $\tau_{\text{target}}^{\text{PaySim}} = \{R_3 = 0, \text{int} = 1\}$
- Rule $\text{isFraud}(X) \leftarrow \text{NOT}(\text{isNormal}(X) > -0.807(X_0), \text{pred}(X_0))$

Performance	∂ILP	DT	DS
Accuracy	0.998	0.999	0.999
Precision	0.973	0.971	0.984
Recall	0.981	0.965	0.981
F1	0.982	0.969	0.984
MCC	0.688	0.693	0.702

Pipeline



Program Template

- Rules (τ_p, τ_q) define predicate p template
- $\tau = (R_3, \text{int})$ defines the range of clauses C to generate, each clause consists of two atoms
- \exists Number of existential predicates
- int A flag to use an intensional (auxiliary) predicate in generated clauses
- rules are defined for each predicate P
- arity_P The arity of an auxiliary predicate P_q
- Generates a set of clauses:

$$\begin{aligned} &C_1^{(1)}, C_1^{(2)}, C_1^{(3)}, \dots, C_1^{(P_1)}, C_1^{(P_1+1)} \\ &C_2^{(1)}, C_2^{(2)}, C_2^{(3)}, \dots, C_2^{(P_2)}, C_2^{(P_2+1)} \\ &C_3^{(1)}, C_3^{(2)}, C_3^{(3)}, \dots, C_3^{(P_3)}, C_3^{(P_3+1)} \end{aligned}$$

Induction as Satisfiability

For each P , ∂ILP learns a weight matrix W_P , to find a set of clauses best explaining Positive, Negative instances of a Target predicate

Inference Steps

$\text{edge}(x,y) \leftarrow \text{connected}(X,Y) = \text{edge}(X,Y)$
 $\text{edge}(x,c) \leftarrow \text{connected}(X,Y) = \text{edge}(X,Y), \text{connected}(Z,Y)$
 $\text{edge}(c,x)$
 $C_{X1} = \{\text{edge}(x,y), \text{edge}(y,c), \text{edge}(c,x)\}$
 $C_{X2} = C_{X1} \cup \{\text{connected}(x,y), \text{connected}(y,c), \text{connected}(c,x)\}$
 $C_{X3} = C_{X2} \cup \{\text{connected}(x,y), \text{connected}(y,c), \text{connected}(c,x)\}$
 $C_{X4} = C_{X3} \cup \{\text{connected}(x,y), \text{connected}(y,c), \text{connected}(c,x)\}$

Pros:

- ∂ILP can generalize from a small amount of data; not data-hungry
- Creates recursion predicates
- Provides shorter explainable rules than Decision Tree

Cons:

- Hard to scale
- Did not outperform other techniques
- Required to define Program Template, and to convert dataset to binary dataset



Amsterdam
Data Science